Policy wording extraction

AI service: policy knowledge · Run 09 Nov 2026, 02:01 UTC · **All targets met**

What is measured

Whether every clause of a policy wording is found, and every limit, percentage, and waiting period read exactly, with its exceptions, cross-references, and footnotes. This is what a policy manager then checks before a second manager publishes the version.

Data

Three synthetic policy wordings with known answers, including a benefits table crossing a page break, merged cells, footnotes, exceptions inside exclusions, and a clause split across pages.

Results

Measure	Result	Target	
Critical value accuracy before human review	**49 / 49 (100.0%)**	≥ 98%	**Met**
Clauses found	**53 / 53 (100.0%)**	≥ 95%	**Met**
Values marked present without a source quote	**0**	0	**Met**
Exceptions captured	**8 / 8 (100.0%)**		
Cross-references found	**11 / 11 (100.0%)**		
Footnotes attached	**2 / 2 (100.0%)**		
Policies that failed	**0**	0	**Met**
Cost for the three wordings	**USD 0.0136**		

Found or read correctly (%)

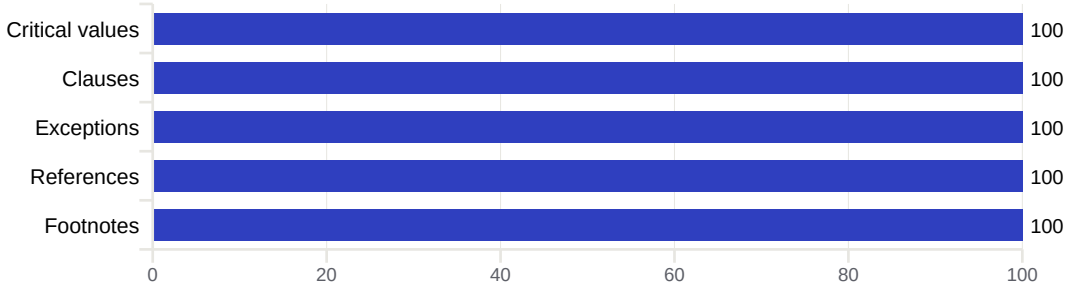

Observations

- Model deepseek/deepseek-v4.1-flash-1017. Pipeline policy-1.
- No errors.

Limitations

- Three wordings written for testing. Real wordings are longer and less regular.

Reproduce

```
cd services/ai; python -m uv run python evals/run_policy_extraction.py
```