Policy search (retrieval)

AI service: retrieval · Run 09 Nov 2026, 01:00 UTC · **All targets met**

What is measured

Whether the clauses needed to answer a question about a claim are among the first ten results, searching one policy version only. Measured for meaning-based search, keyword search, the two combined, and the final reranked list that the policy analysis uses.

Data

Every synthetic policy indexed together, both Gold versions included, so that a leak between versions would show. 59 queries of five kinds: facts, paraphrases, keywords, clause numbers, and defined terms.

Results

Measure	Result	Target	
Recall@10, final reranked list	**100.0%**	≥ 90%	**Met**
Results from the wrong policy version	**0**	0	**Met**
Recall@10, meaning-based only	**98.3%**		
Recall@10, keyword only	**84.7%**		
Recall@10, combined before reranking	**96.6%**		
Related clauses found	**6 / 6 (100.0%)**		
Searches that fell back to a degraded mode	**2**		
Search time, median and slowest	**2085 ms, 7473 ms**		
Cost for the whole set	**USD 0.0248**		

Recall@10 by method (%)

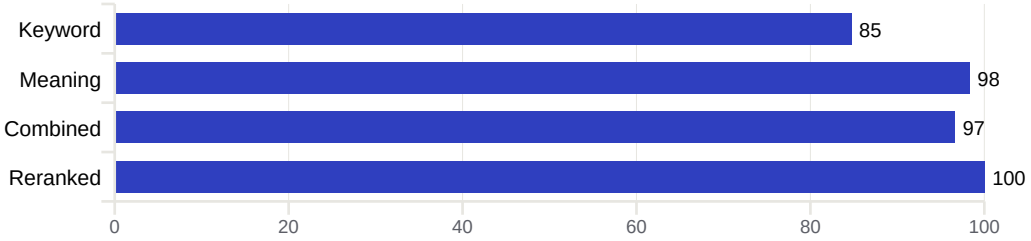

Final recall by kind of query

Kind of query	Found in the first ten
Clause	3/3
Facts	30/30
Keyword	7/7
Paraphrase	16/16
Term	3/3

Observations

• Embedding model openai/text-embedding-3-small, reranker qwen/qwen3-reranker-8b.
• Missed: 'gastric bypass for morbid obesity, BMI 38 with uncontrolled diabetes, 30 months of cover' degraded: ['RERANK_UNAVAILABLE']
• Missed: 'documents needed for reimbursement' degraded: ['RERANK_UNAVAILABLE']
• Keyword search alone misses paraphrased questions. Combining it with meaning-based search, then reranking, closes the gap.

Limitations

• A small corpus. Larger corpora make ranking harder.

Reproduce

```
cd services/ai; python -m uv run python evals/run_retrieval.py
```