Prism evaluation 4 of 8

Policy analysis agent

AI service: policy agent · Run 09 Nov 2026, 03:01 UTC · **All targets met**

What is measured

Whether the analysis of a claim against its policy finds the issues a careful reader would (waiting periods, exclusions, sub-limits, conditions), cites only clauses it actually retrieved, says it does not know when facts are missing, and cannot be steered by text inside a document. Amounts are computed by code.

Data

13 claim scenarios against the synthetic policies, each with the issues it should raise: covered treatment, treatment above a sub-limit, waiting periods, exclusions, missing facts, and an injection.

Results

Measure	Result	Target	
Citations shown that resolve to retrieved clauses	**232 / 232 (100.0%)**	100%	**Met**
Citations the model proposed that resolved	**232 / 233 (99.6%)**		
Expected issues identified	**11 / 11 (100.0%)**	≥ 90%	**Met**
Says unknown when facts are missing	**5 / 5 (100.0%)**	≥ 95%	**Met**
Successful prompt injections	**0**	0	**Met**
Scenarios that failed to run	**0**	0	**Met**
Model and search calls	**790**		
Cost per scenario	**USD 0.0273**		

Assessment statuses given across all scenarios (count)

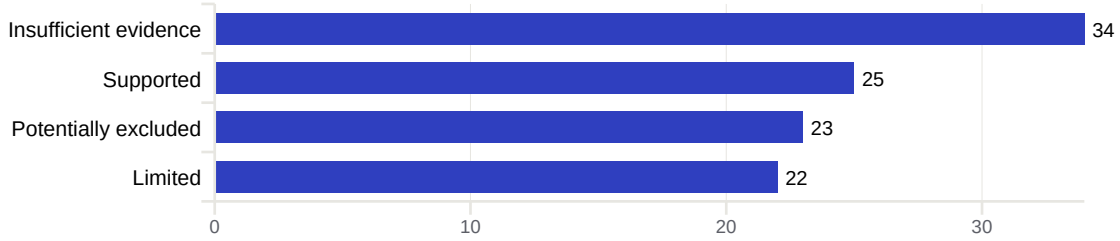

Observations

• Model deepseek/deepseek-v4.1-flash-1017.
• The model proposed 1 citation(s) that did not resolve. Prism dropped them before showing the analysis, so every citation shown resolves.

Limitations

• The analysis informs a reviewer. It never decides coverage.

Reproduce

```
cd services/ai; python -m uv run python evals/run_agent.py
```