Risk summary (language model)

AI service: risk summary · Run 09 Nov 2026, 01:01 UTC · **All targets met**

What is measured

Whether the written summary of a claim's risk signals is grounded and never accusatory. Code checks every summary: each number must appear in the signals, each citation must name a real signal, and words such as fraud or fraudulent are refused. A summary that fails is repaired once, or not shown.

Data

The signals of the severe and disguised demonstration claims, exactly as core sends them, summarized 10 times with the configured model.

Results

Measure	Result	Target	
Summaries that passed every check	**9 / 10 (90.0%)**		
Summaries shown that fail a check	**0**	0, enforced by code	**Met**
Attempts repaired or refused before passing	**9**		
Cost per summary	**USD 0.0006**		

Observations

• Example, severe claim: “Several high-severity document inconsistencies need attention first: the patient's name differs across documents [S1], and two bills are dated before the admission — the hospital bill 5 Jul 2026 [S3] and the itemized invoice 6 Aug 2026 [S4]. The itemized invoice's line items add up to INR 3,10,000.00 against a stated total of INR 3,80,000.00 [S5], and the claim asks for INR 7,40,000.00 while the bills total INR 4,92,600.00 [S7] and the sum insured is INR 5,00,000.00 [S8]. Duplicate indicators also appear: the hospital bill matches a document in another claim [S10] and an invoice number is reused at the same hospital [S11]. Statistically, the claim is unusual compared with the baseline [S12][S13]. The claim goes to senior review because of these signals, and the reviewer decides.”
• Example, disguised claim: “The highest-priority item is a duplicate indicator: 5 of 7 line items on the hospital bill match, by description and amount, a bill from the same hospital in another claim [S1]. A similar claim for the same patient at the same hospital, 18 days apart, is also noted [S2]. Separately, the claim is statistically unusual compared with the synthetic baseline: Isolation Forest scores it as more unusual than 99.84% of baseline claims [S3], and Local Outlier Factor scores it as more unusual than 99.88% [S4]. Because of these signals, the claim goes to senior review, and the reviewer decides.”

Limitations

• Passing the checks means grounded and non-accusatory, not the best possible wording. The examples are here to be read.

Reproduce

```
cd services/ai; python -m uv run python evals/run_risk_summary.py
```