Risk summary (language model)

AI service: risk summary · Run 09 Nov 2026, 02:13 UTC · **All targets met**

What is measured

Whether the written summary of a claim's risk signals is grounded and never accusatory. Code checks every summary: each number must appear in the signals, each citation must name a real signal, and words such as fraud or fraudulent are refused. A summary that fails is repaired once, or not shown.

Data

The signals of the severe and disguised demonstration claims, exactly as core sends them, summarized 20 times with the configured model.

Results

Measure	Result	Target	
Summaries that passed every check	**20 / 20 (100.0%)**		
Summaries shown that fail a check	**0**	0, enforced by code	**Met**
Attempts repaired or refused before passing	**3**		
Cost per summary	**USD 0.0002**		

Observations

• Example, severe claim: “The reviewer's attention should go first to the duplicate-document indicators: the hospital bill is the same file as a document in another claim [S10], and the itemized invoice reuses an invoice number from a bill at the same hospital in another claim [S11]. Document consistency also fails: the patient's name differs across documents [S1], and both the hospital bill dated 5 Jul 2026 [S3] and the itemized invoice dated 6 Aug 2026 [S4] predate the admission on 10 Aug 2026. Financial checks show the itemized invoice total of INR 3,80,000.00 against line items adding to INR 3,10,000.00 [S5], a claim of INR 7,40,000.00 against bills totaling INR 4,92,600.00 [S7] and above the sum insured of INR 5,00,000.00 [S8]. Statistically, the claim is unusual compared with the synthetic baseline [S12] [S13]. The claim goes to senior review because of these signals, and the reviewer decides.”
• Example, disguised claim: “The reviewer's attention should go first to [S1], a high-severity rule signal: 5 of 7 line items on the hospital bill match, by description and amount, a bill from the same hospital in another claim. [S2] adds that another claim for the same patient at the same hospital, 18 days apart, is for a similar amount. Separately, two statistical detectors mark the claim as statistically unusual compared with the synthetic baseline: [S3] (more unusual than 99.84% of baseline claims) and [S4] (more unusual than 99.88%). The claim goes to senior review because of these signals, and the reviewer decides.”

Limitations

• Passing the checks means grounded and non-accusatory, not the best possible wording. The examples are here to be read.

Reproduce

```
cd services/ai; python -m uv run python evals/run_risk_summary.py
```